

```
In [4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
In [5]: from sklearn.datasets import load_boston
boston = load_boston()
```

C:\Users\Joshua Deshmukh\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load_boston is deprecated; `load_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch_california_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

```
In [6]: boston.DESCR
```

```
Out[6]: ".. _boston_dataset:\n\nBoston house prices dataset\n-----\n\n**Data Set Characteristics:** \n\n :Number of Instances: 506 \n\n :Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.\n\n :Attribute Information (in order):\n      - CRIM      per capita crime rate by town\n      - ZN        proportion of residential land zoned for lots over 25,000 sq.ft.\n      - INDUS     proportion of non-retail business acres per town\n      - CHAS      Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)\n      - NOX       nitric oxides concentration (parts per 10 million)\n      - RM        average number of rooms per dwelling\n      - AGE       proportion of owner-occupied units built prior to 1940\n      - DIS       weighted distances to five Boston employment centres\n      - RAD       index of accessibility to radial highways\n      - TAX       full-value property-tax rate per $10,000\n      - PTRATIO   pupil-teacher ratio by town\n      - B         1000(Bk - 0.63)^2 where Bk is the proportion of black people by town\n      - LSTAT     % lower status of the population\n      - MEDV      Median value of owner-occupied homes in $1000's\n\n :Missing Attribute Values: None\n\n :Creator: Harrison, D. and Rubinfeld, D.L.\n\nThis is a copy of UCI ML housing dataset.\nhttps://archive.ics.uci.edu/ml/machine-learning-databases/housing/\n\n\nThis dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.\n\nThe Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic\nprices and the demand for clean air', J. Environ. Economics & Management,\nvol.5, 81-102, 1978.  Used in Belsley, Kuh & Welsch, 'Regression diagnostics\n...', Wiley, 1980.  N.B. Various transformations are used in the table on\npages 244-261 of the latter.\n\nThe Boston house-price data has been used in many machine learning papers that address regression\nproblems. \n\n.. topic:: References\n\n - Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.\n - Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.\n"
```

```
In [7]: df = pd.DataFrame(boston.data, columns=boston.feature_names)
df['Price'] = boston.target
```

In [8]:

df

Out[8]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	Price
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33	36.2
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67	22.4
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88	11.9

506 rows × 14 columns

In [9]:

df.describe()

Out[9]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	12.659589
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	7.145456
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	1.730000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	6.950000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	11.360000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	16.950000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	37.970000

```
In [10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype  
---  -
0   CRIM         506 non-null   float64
1   ZN           506 non-null   float64
2   INDUS        506 non-null   float64
3   CHAS         506 non-null   float64
4   NOX          506 non-null   float64
5   RM           506 non-null   float64
6   AGE          506 non-null   float64
7   DIS          506 non-null   float64
8   RAD          506 non-null   float64
9   TAX          506 non-null   float64
10  PTRATIO      506 non-null   float64
11  B            506 non-null   float64
12  LSTAT        506 non-null   float64
13  Price        506 non-null   float64
dtypes: float64(14)
memory usage: 55.5 KB
```

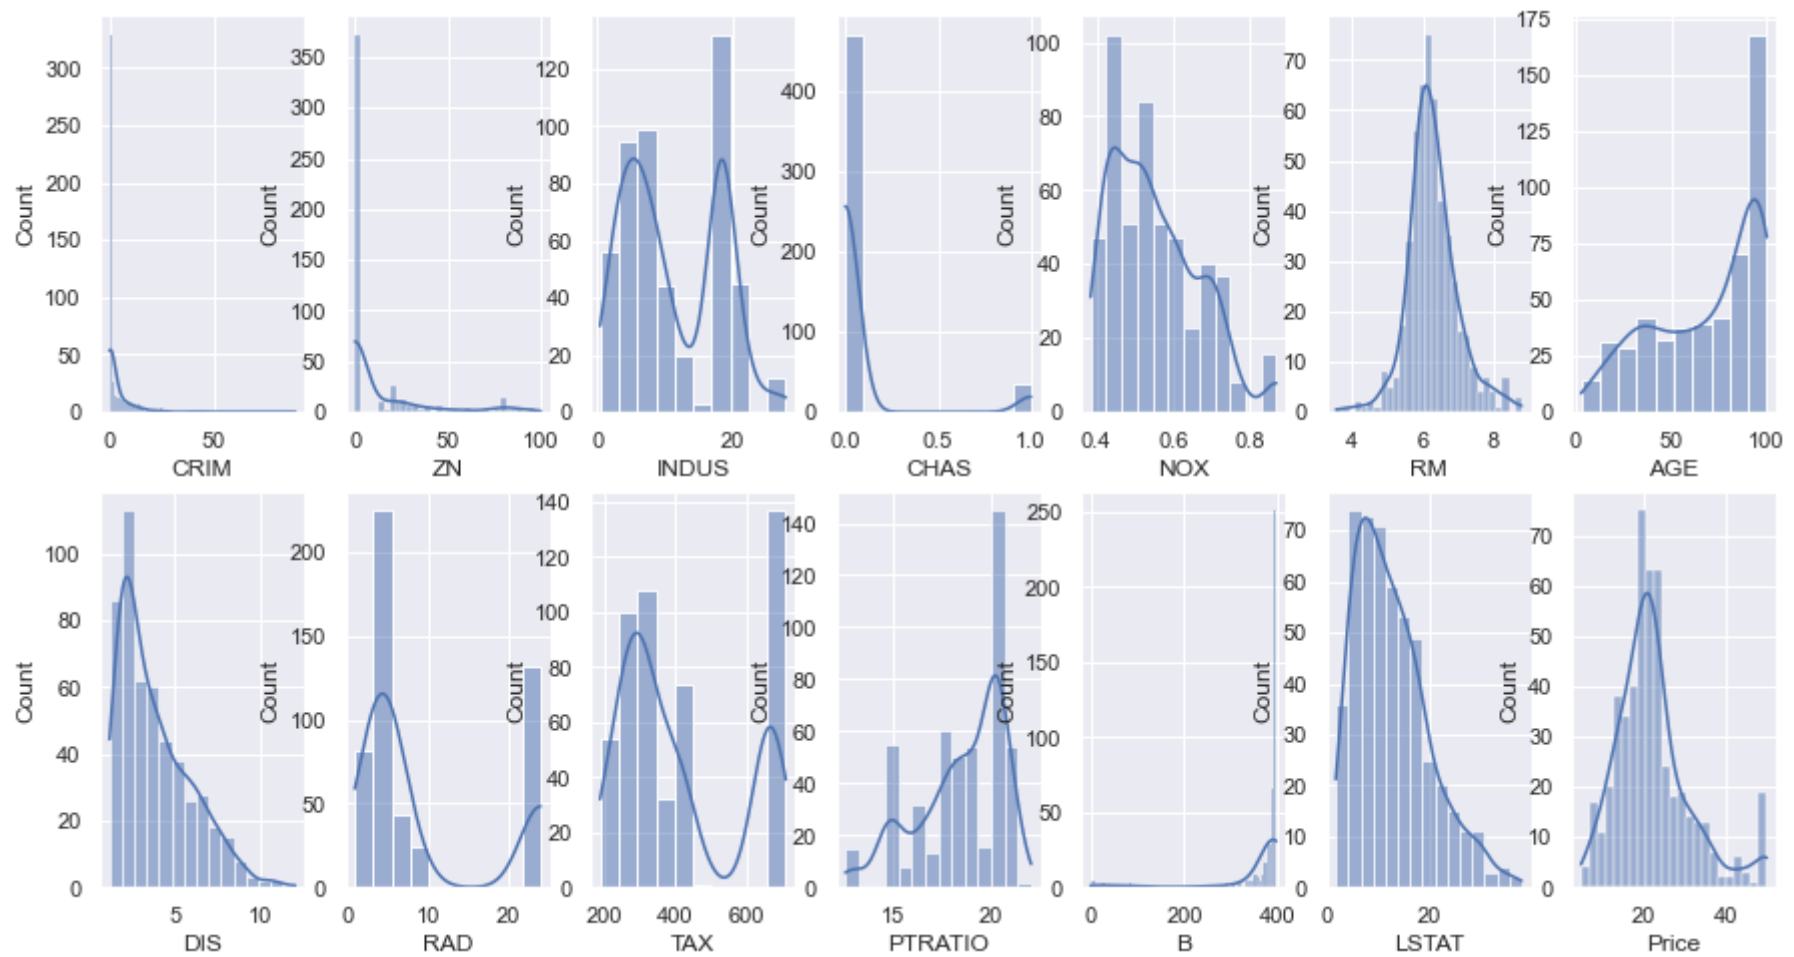
```
In [11]: df.isnull().sum()
```

```
Out[11]: CRIM      0
ZN            0
INDUS         0
CHAS          0
NOX           0
RM            0
AGE           0
DIS           0
RAD           0
TAX           0
PTRATIO       0
B             0
LSTAT         0
Price         0
dtype: int64
```

```


In [13]: row = 2
col = 7
s = 0
fig, ax=plt.subplots(nrows=row, ncols=col, figsize=(15,8))
for i in range(row):
    for j in range(col):
        sns.histplot(x=df.columns[s], data=df, kde=True, ax=ax[i][j])
        s+=1

```



```
In [14]: fig, ax = plt.subplots(figsize = (18,12))  
sns.heatmap(data=df.corr(), annot=True, ax=ax)
```

Out[14]: <AxesSubplot:>




```
In [15]: cor=df.corr()
features=[]
for i in cor['Price'].index:
    if abs(cor['Price'][i]) < 0.45:
        features.append(i)
df = df.drop(columns=features)
```

```
In [16]: df
```

```
Out[16]:
```

	INDUS	RM	TAX	PTRATIO	LSTAT	Price
0	2.31	6.575	296.0	15.3	4.98	24.0
1	7.07	6.421	242.0	17.8	9.14	21.6
2	7.07	7.185	242.0	17.8	4.03	34.7
3	2.18	6.998	222.0	18.7	2.94	33.4
4	2.18	7.147	222.0	18.7	5.33	36.2
...
501	11.93	6.593	273.0	21.0	9.67	22.4
502	11.93	6.120	273.0	21.0	9.08	20.6
503	11.93	6.976	273.0	21.0	5.64	23.9
504	11.93	6.794	273.0	21.0	6.48	22.0
505	11.93	6.030	273.0	21.0	7.88	11.9

506 rows × 6 columns

```
In [17]: x = df.drop(labels=['Price'],axis=1)
y = df['Price']
```

```
In [18]: x.head()
```

```
Out[18]:
```

	INDUS	RM	TAX	PTRATIO	LSTAT
0	2.31	6.575	296.0	15.3	4.98
1	7.07	6.421	242.0	17.8	9.14
2	7.07	7.185	242.0	17.8	4.03
3	2.18	6.998	222.0	18.7	2.94
4	2.18	7.147	222.0	18.7	5.33

```
In [19]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.3, random_state=42)
```

```
In [20]: x_train.shape , x_test.shape
```

```
Out[20]: ((354, 5), (152, 5))
```

```
In [21]: from sklearn.linear_model import LinearRegression
lr = LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[21]: LinearRegression()
```

```
In [22]: lr.coef_
```

```
Out[22]: array([ 8.89192908e-02,  4.58472965e+00, -3.21801353e-03, -8.30469649e-01,
                -6.14838082e-01])
```

```
In [23]: lr.intercept_
```

```
Out[23]: 17.150526471659887
```

```
In [24]: y_pred = lr.predict(x_test)
y_pred
```

```
Out[24]: array([ 26.62981059,  31.10008241,  16.95701338,  25.59771173,  18.09307064,
  22.90871478,  17.61601889,  13.53637406,  20.48659658,  19.63042973,
  20.14356915,  21.23936603, -2.17703728,  22.46574239,  19.55741359,
  24.67368263,  19.10912156,   3.38842032,  38.69074984,  17.05082016,
  26.1619823 ,  27.58438964,  11.74375654,  24.17459161,  17.7242101 ,
  13.56944671,  22.81648456,  19.18372228,  18.46079157,  18.70360048,
  19.2864636 ,  25.64061437,  25.10598582,  18.15940204,  14.25506589,
  21.85273114,  32.68155189,  20.99806076,  20.54233085,  25.18099014,
  12.5229096 ,  28.12041259,  39.52307563,  18.48292269,  25.8060045 ,
  15.51396261,  14.30180164,  26.47373087,  18.0995971 ,  30.77991619,
  23.61711953,  33.50182842,  16.09617732,  25.6923864 ,  38.16291719,
  22.27510835,  18.15355162,  30.2656657 ,  24.76833162,  15.09517744,
  25.49349461,  31.94366429,  29.9910815 ,  17.2733645 ,  27.56807723,
  12.05824047,  18.98303931,  25.66074343,  28.98449592,  15.55155416,
  19.97704179,  26.02936009,  12.05329111,  21.7655206 ,  23.31666209,
   6.03391556,  20.01221693,  38.12144014,  16.8875505 ,  10.72075752,
  22.46088278,   9.36198278,  24.09343534,   7.14093828,  22.17983415,
  28.15116909,  20.46791378,  26.25420747,  27.1620727 ,  20.70409881,
  24.1563809 ,   7.77940126,  21.54982615,  20.04604896,  11.78985685,
  22.46163913,  22.48055429, -0.07257436,  18.36539259,  17.31496289,
  21.06156734,  24.80376727,   8.84307115,  20.98617093,  25.51248652,
  13.50230783,  18.76790985,  26.96699475,  22.8688555 ,  27.5586772 ,
  11.78407124,  19.09435201,  24.89518268,  24.32419065,  31.1937723 ,
  18.97618967,  32.91808139,  14.96546128,  19.58974059,  28.640961 ,
  18.70459566,  26.86127531,  14.47455498,  23.13989195,  26.86074285,
  23.56807441,  27.48834026,  31.80605265,  24.3571917 ,  37.91833369,
  11.93795084,  26.54212343,  19.26752655,  19.29065717,  11.75099637,
  21.90166758,  22.66230052,  32.72490364,  31.18752369,  16.58461442,
  17.88939265,  29.94107039,  23.23104109,  11.78611835,   8.36515471,
  24.64330163,  24.19025911,  17.20942048,  13.87458909,  39.4187899 ,
  19.37592804,  18.39248868])
```

```
In [25]: pred_df=pd.DataFrame(np.c_[ y_test, y_pred], columns = ["Price_original","Price_predicted"])
pred_df
```

```
Out[25]:
```

	Price_original	Price_predicted
0	23.6	26.629811
1	32.4	31.100082
2	13.6	16.957013
3	22.8	25.597712
4	16.1	18.093071
...
147	17.1	17.209420
148	14.5	13.874589
149	50.0	39.418790
150	14.3	19.375928
151	12.6	18.392489

152 rows × 2 columns

```
In [26]: from sklearn.metrics import r2_score
lr.score(x_test,y_test)
```

```
Out[26]: 0.6499135956539925
```

```
In [27]: from sklearn.ensemble import RandomForestRegressor
```

```
In [28]: rf = RandomForestRegressor()
rf.fit(x_train,y_train)
```

```
Out[28]: RandomForestRegressor()
```

```
In [29]: y_pred = rf.predict(x_test)
y_pred
```

```
Out[29]: array([22.79 , 31.42 , 17.212, 23.11 , 15.309, 21.744, 20.294, 14.599,
                21.656, 21.527, 21.135, 20.206, 12.55 , 21.966, 18.41 , 25.179,
                19.719,  8.576, 46.466, 15.646, 24.447, 24.189, 14.054, 23.025,
                14.996, 14.878, 24.305, 16.193, 20.784, 21.237, 18.798, 23.29 ,
                26.444, 21.41 , 10.588, 16.656, 35.483, 19.321, 20.404, 23.269,
                15.53 , 29.268, 45.572, 20.356, 23.262, 14.475, 17.28 , 23.639,
                15.239, 29.454, 22.301, 34.702, 17.868, 25.872, 43.895, 21.051,
                15.809, 33.086, 22.272, 19.617, 25.793, 34.747, 29.585, 19.324,
                26.71 , 20.405, 15.626, 22.965, 28.081, 20.448, 20.465, 31.97 ,
                10.85 , 22.128, 22.047,  7.744, 20.521, 47.725, 12.196, 11.4 ,
                22.783,  8.313, 23.154,  9.293, 21.6 , 27.652, 15. , 23.166,
                23.439, 17.367, 21.153,  8.465, 19.754, 19.614, 32.306, 19.631,
                26.371, 10.978, 14.92 , 12.454, 19.932, 27.169, 11.098, 21.447,
                21.883, 12.086, 18.351, 24.556, 20.315, 23.644,  7.698, 13.563,
                23.387, 25.076, 32.805, 15.947, 43.107, 16.964, 17.706, 23.838,
                20.121, 24.059,  9.675, 20.379, 23.965, 21.377, 24.495, 35.854,
                18.997, 47.512, 17.886, 22.144, 19.358, 18.827, 14.21 , 21.555,
                21.367, 32.007, 29.392, 16.872, 17.313, 25.865, 20.816, 20.508,
                8.312, 22.072, 18.461, 12.582, 13.678, 42.133, 15.365, 15.653])
```

```
In [30]: rf.score(x_test,y_test)
```

```
Out[30]: 0.8013549394401752
```

```
In [ ]:
```