

```
In [3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
```

```
In [4]: df=pd.read_csv('Social_Network_Ads.csv')
```

```
In [5]: df
```

```
Out[5]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...	...	...	...	...	...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
In [6]: df.isnull().sum()
```

```
Out[6]: User ID      0
Gender      0
Age         0
EstimatedSalary  0
Purchased   0
dtype: int64
```

```
In [14]: df1=df.drop(['User ID'],axis=1)
```

```
In [15]: df1.head(10)
```

```
Out[15]:
```

	Gender	Age	EstimatedSalary	Purchased
0	Male	19	19000	0
1	Male	35	20000	0
2	Female	26	43000	0
3	Female	27	57000	0
4	Male	19	76000	0
5	Male	27	58000	0
6	Female	27	84000	0
7	Female	32	150000	1
8	Male	25	33000	0
9	Female	35	65000	0

```
In [17]: from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression
```

```
In [18]: enc=LabelEncoder()
```

```
In [21]: df1['Gender']=enc.fit_transform(df1['Gender'])
```

```
In [23]: df1.head()
```

```
Out[23]:
```

	Gender	Age	EstimatedSalary	Purchased
0	1	19	19000	0
1	1	35	20000	0
2	0	26	43000	0
3	0	27	57000	0
4	1	19	76000	0

```
In [25]: x=df1.drop(labels=['Purchased'],axis=1)
```

```
In [26]: y=df1['Purchased']
```

```
In [27]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=10)
```

```
In [28]: x_train.shape, x_test.shape
```

```
Out[28]: ((300, 3), (100, 3))
```

```
In [29]: logr = LogisticRegression()  
logr.fit(x_train,y_train)
```

```
Out[29]: LogisticRegression()
```

```
In [30]: logr.score(x_test,y_test)
```

```
Out[30]: 0.69
```


```
In [31]: y_pred = logr.predict(x_test)
```

```
In [32]: y_pred
```

```
Out[32]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)
```

```
In [33]: sns.countplot(x='Purchased',data=df1)
```

```
Out[33]: <AxesSubplot:xlabel='Purchased', ylabel='count'>
```



```
In [34]: from sklearn.metrics import confusion_matrix
```

```
In [35]: confusion_matrix(y_test,y_pred)
```

```
Out[35]: array([[69,  0],
                [31,  0]], dtype=int64)
```

```
In [36]: from sklearn.metrics import classification_report as cr
print(cr(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.69	1.00	0.82	69
1	0.00	0.00	0.00	31
accuracy			0.69	100
macro avg	0.34	0.50	0.41	100
weighted avg	0.48	0.69	0.56	100

C:\Users\Joshua Deshmukh\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Joshua Deshmukh\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Joshua Deshmukh\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

```
In [38]: from statsmodels.stats.outliers_influence import variance_inflation_factor
variables = df[['Gender', 'EstimatedSalary', 'Age']]
vif = pd.DataFrame()
vif["VIF"] = [variance_inflation_factor(variables.values, i) for i in range(variables.shape[1])]
vif["features"] = variables.columns
vif
```

```
Out[38]:
```

	VIF	features
0	1.783132	Gender
1	4.601775	EstimatedSalary
2	5.122407	Age

In [ ]: